Kernel Latent SVM for Visual Recognition: Reproduction and Analysis

Pranay Vishwakarma
Roll No: 230083
Advanced Machine Learning
Mid-Semester Examination — Part B

I. Paper Summary

The paper proposes **Kernel Latent SVM (KLSVM)**, a framework that unifies latent variable modeling with kernel methods for visual recognition. Standard latent SVMs (LSVM) use linear scoring functions

$$f_w(x) = \max_h w^\top \phi(x, h) \tag{1}$$

which limits their representational capacity. KLSVM reformulates the optimization in the dual space, replacing dot-products with kernel functions to obtain nonlinear classifiers while retaining latent structure.

The iterative algorithm alternates between latent variable inference and kernel SVM training. The paper reports improvements on multiple benchmarks including mammal classification (84.49% vs 75.07%), CIFAR-10 (47.72% vs 39.42%), and object interaction recognition (66.42% vs 46.33%).

II. Reproduction Setup and Results

We reproduced Section 4.2 (latent subcategories on CIFAR-10) using a 4-class subset (airplane, automobile, bird, cat). The setup included 200 training and 100 test samples per class, HOG features, Histogram Intersection Kernel (HIK), regularization parameter $C = 0.01$, and $K = 3$ latent subcategories.

Fig. 1. Sample images from the CIFAR-10 subset used in our experiments.

Our ordering (Linear SVM > Kernel SVM > KLSVM ≈ LSVM) reverses the paper's finding because our training data is much smaller than the original dataset (200 vs 10,000 samples per class).

The subcategory expansion increases the feature dimension from 288 to 864, creating a severely underdetermined problem. The alternating optimization fails to converge, with approximately 600 subcategory labels changing per iteration. This indicates that KLSVM requires large datasets to discover meaningful latent substructures.

TABLE I
Comparison with Reported Paper Results

Method	Our Accuracy	Paper Accuracy
Linear SVM	**49.25%**	38.32%
Kernel SVM (HIK)	47.75%	44.04%
Linear LSVM	39.00%	39.42%
KLSVM	39.25%	**47.72%**

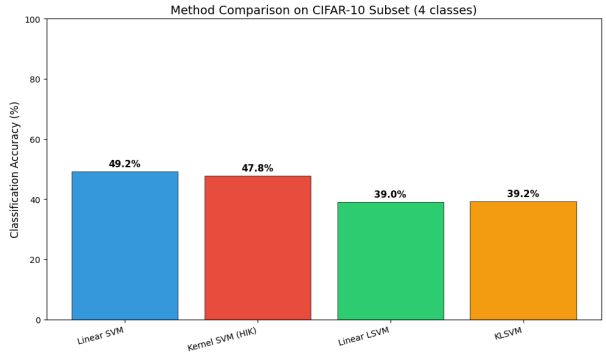

Fig. 2. Accuracy comparison across evaluated methods.

III. Ablation Findings

Ablation 1 (Removing subcategories): Performance improved from 39.25% to 47.75% (+8.5%). This shows that latent subcategories fragment the dataset when training data is limited.

Ablation 2 (Replacing HIK with linear kernel): Performance remains similar (39.00% vs 39.25%), indicating that kernel choice is not the bottleneck.

These results show that KLSVM's advantages are strongly *data-dependent*. With insufficient data, increased model complexity harms performance.

IV. Failure Mode

Cat vs. dog classification with Gaussian noise reveals a consistent performance gap where KLSVM underperforms standard SVM by approximately 3–6%.

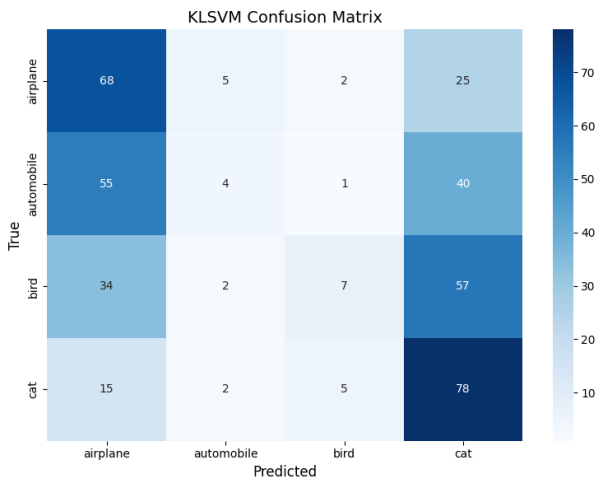

Fig. 3. Confusion matrix for the KLSVM classifier.

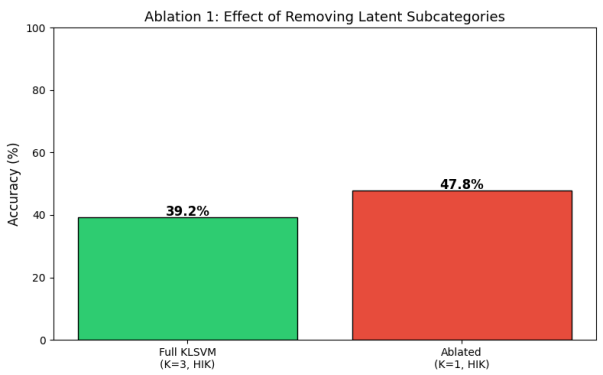

Fig. 4. Effect of removing latent subcategories.

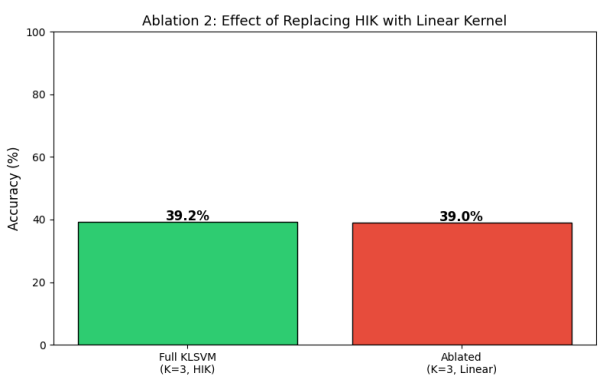

Fig. 5. Effect of replacing HIK with a linear kernel.

Two assumptions of the model are violated:

- **Optimization Landscape Assumption:** The alternating optimization used in latent SVM is non-convex and may converge to poor local minima. When classes are visually similar, k-means initialization generates arbitrary subcategories that reflect noise patterns rather than meaningful structure.
- **Feature Discriminativeness Assumption:** The joint feature representation $\phi(x, h)$ must remain discriminative across latent configurations. Cat and dog images produce nearly identical HOG descriptors, making subcategory-based separation ineffective.

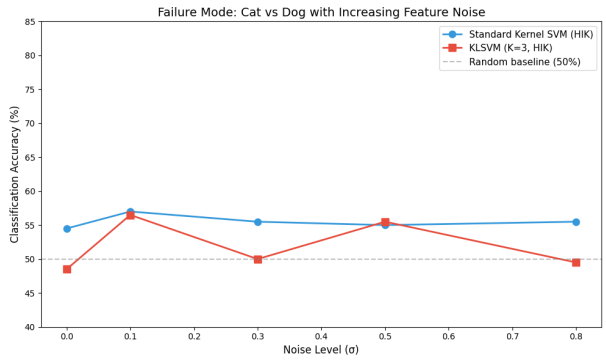

Fig. 6. Failure mode experiment: cat vs. dog classification with increasing noise.

A potential solution is to automatically determine the number of latent subcategories using silhouette scores or gap statistics and default to $K = 1$ when no meaningful structure exists.

V. Reflection

What surprised me: The sensitivity of KLSVM to dataset size was greater than expected. The paper's results rely heavily on very large training sets.

Future work: I would revisit the full coordinate ascent optimization (Eq. 8–9) using larger datasets and explore composite kernels for structured latent variables on the VOC3000 dataset.

References

1) J. Hu et al., "Kernel Latent SVM for Visual Recognition," in *Advances in Neural Information Processing Systems (NeurIPS)*, 2012.
2) P. Felzenszwalb et al., "Object Detection with Discriminatively Trained Part-Based Models," *IEEE TPAMI*, 2010.
3) S. Maji et al., "Classification Using Intersection Kernel SVMs is Efficient," in *CVPR*, 2008.
4) C. Yu and T. Joachims, "Learning Structural SVMs with Latent Variables," in *ICML*, 2009.
5) A. Krizhevsky, "Learning Multiple Layers of Features from Tiny Images," 2009.